

Final_RL

May 12, 2022

0.1 Alli Khadga Jyoth - 19024 DSE

0.2 Akash Kumar Singh - 19023 DSE

```
[1]: from gym import Env
import gym
import pygame
from gym.spaces import Discrete, Box, Dict
import numpy as np
import random
import matplotlib.pyplot as plt
from IPython.display import clear_output
np.set_printoptions(linewidth=500,threshold=np.inf)
```

0.3 ATC Environment

```
[2]: class AirTraffic():
    def __init__(self, planes: int = 2, grid_size: list = [5, 5], radius: int = 2, seed = None, destinations: int = None ):
        self.radius = radius
        self.planes = planes
        self.destinations = destinations
        self.GRID_DIM = grid_size
        self.seed = seed
        self._max_dist = np.sqrt(self.GRID_DIM[0]**2 + self.GRID_DIM[1]**2)
        self._action_to_direction = {
            0: np.array([-1, 0]),
            1: np.array([-1, 1]),
            2: np.array([0, 1]),
            3: np.array([-1, -1]),
            4: np.array([0, -1]),
        }

        self._ACTIONLOOKUP = {
            0: 'Up',
            1: 'front right',
            2: 'right',
```

```

        3: 'front left',
        4: 'left'
    }
    self.GRID_DIM = np.array(self.GRID_DIM)
    self.GRID = np.zeros(self.GRID_DIM)
    self.action_space = Discrete(len(self._ACTIONLOOKUP.keys()))
    self.state_space = Discrete(self.GRID_DIM[0]*self.GRID_DIM[1])
    self._ob_space = {}
    for p in range(self.planes):
        self._ob_space[f'plane{p}'] = Box(np.array([0,0]), np.array([self.
→GRID_DIM[0]-1,self.GRID_DIM[1] - 1]), shape=(2,), dtype=int)
        self._ob_space[f'dest{p}'] = Box(np.array([0,0]), np.array([self.
→GRID_DIM[0]-1,self.GRID_DIM[1] - 1]), shape=(2,), dtype=int)

    self.observation_space = Dict(self._ob_space)
    self._agent_location = []
    self._target_location = []
    # Randomly initialize the start points of planes and destinations
    if self.seed is not None:
        np.random.seed(seed = seed)
    for p in range(self.planes):
        while (pc:= [np.random.randint(self.GRID_DIM[0]),np.random.
→randint(self.GRID_DIM[1])]) in self._agent_location:
            continue
        self._agent_location.append(pc)
        if self.destinations is None:
            while (dc:= [np.random.randint(self.GRID_DIM[0]),np.random.
→randint(self.GRID_DIM[1])]) in self._agent_location + self._target_location:
                continue
            self._target_location.append(dc)
        if self.destinations is not None:
            for d in range(self.destinations):
                while (dc:= [np.random.randint(self.GRID_DIM[0]),np.random.
→randint(self.GRID_DIM[1])]) in self._agent_location + self._target_location:
                    continue
                self._target_location.append(dc)

        extra = np.random.choice(range(len(self._target_location)),self.
→planes-self.destinations)
        for d in extra:
            self._target_location.append(self._target_location[d])
    self._agent_location = np.array(self._agent_location)
    self._target_location = np.array(self._target_location)
    self._count = [0]*self.planes # to check if its already landed/reached_
→destination
    self._done_prev = [False]*self.planes

```

```

self._prev_agent_location = self._agent_location.copy()

def step(self, action):
    self._prev_agent_location = self._agent_location.copy() ##### Find the
    ↪ vector of agent from destination to find ### reward of terminal state for a
    ↪ plane is 0 and that plane cant move now
    moved_plane = [False]*self.planes
    moved_plane = self._move(action)
    done, reward = self._is_over()
    observation = self._get_obs()
    info = self._get_info()
    self._done_prev = done.copy()
    return observation, reward, done, info
def _move(self, action):
    ##### Find the vector of agent from destination to find
    ### reward of terminal state for a plane is _todestination and that
    ↪ plane cant move now
    self._prev_agent_location = self._agent_location
    for plane, act in enumerate(action):
        if np.array_equal(self._agent_location[plane], self.
    ↪ _target_location[plane]):
            self._count[plane] += 1
            continue
            elif np.sign(self._target_location[plane] - self.
    ↪ _prev_agent_location[plane])[0] <= 0:
                self._agent_location[plane] = np.clip(self.
    ↪ _agent_location[plane] + self._action_to_direction[act], [0, 0], [self.
    ↪ GRID_DIM[0] - 1, self.GRID_DIM[1] - 1])
            else:
                self._agent_location[plane] = np.clip(self.
    ↪ _agent_location[plane] + -1*self._action_to_direction[act], [0, 0], [self.
    ↪ GRID_DIM[0] - 1, self.GRID_DIM[1] - 1])

    return self._prev_agent_location == self._agent_location
def _is_over(self):
    done = (self._agent_location == self._target_location).all(axis = 1)
    reward = self._get_reward()
    return done, reward
def _get_reward(self):
    reward = []
    # reward will be intruder + todestination
    distance, closest_dist, _1, _2 = self._get_info()
    for index, plane in enumerate(range(self.planes)):
        if self._count[plane] > 0:
            reward.append(100) ## terminating state reward

```

```

        else:
            self._intruder = 0
            if closest_dist[plane] < self.radius:

                self._intruder = -(self.radius **2 -
→closest_dist[plane]**2)/(self.radius**2/500)
                self._todestination = 100 - distance[plane]
                reward.append(self._todestination + self._intruder)
            return reward

    def _get_obs(self):
        return {'planes': self._agent_location, 'destinations':self.
→_target_location}

    def _get_info(self):
        closest_dist = []
        distance = []
        plane_theta = []
        plane_rho = []
        dne= self._done_prev
        for index,plane in enumerate(range(self.planes)):
            dist = np.linalg.norm([self._agent_location[plane]]*self.
→planes-self._agent_location,axis = 1)
            dist[plane] = np.inf
            dist[dne] = np.inf
            if dne[plane]:
                dist += np.inf
                plane_rho.append(np.inf)
                plane_theta.append(np.inf)
                closest_dist .append(np.min(dist))
            else:
                closest_dist .append(np.min(dist))
                arg_plane = np.argmin(dist)

                if (self._prev_agent_location==None).all():
                    self._prev_agent_location = self._agent_location.copy()
                    ownship = self._agent_location[plane]-self.
→_prev_agent_location[plane]
                    intruder = self._agent_location[arg_plane]-self.
→_prev_agent_location[arg_plane]
                    closest_path = self._agent_location[arg_plane]-self.
→_agent_location[plane]
                    y = np.array([intruder[1],ownship[1]])
                    x = np.array([intruder[0],ownship[0]])
                    degree = np.rad2deg(np.arctan2(y,x))
                    deg = degree[1]-degree[0]

```

```

        alpha = np.rad2deg(np.arctan2(ownship[1],ownship[1]))
        beta = np.rad2deg(np.
→arctan2(closest_path[1],closest_path[0]))
        deg_rho = (alpha-beta)
        plane_rho.append(deg_rho)
        plane_theta.append( deg)
        distance .append(np.linalg.norm(self.
→_target_location[plane]-self._agent_location[plane]))
        return distance,closest_dist,plane_theta,plane_rho

def _did_collide(self):
    distance,closest_dist,plane_theta,plane_rho = self._get_info()
    collide = np.array(closest_dist).all()
    return collide

def _state_to_index(self):
    _,closest_dist,plane_theta,plane_rho = self._get_info()
    dis_bin = np.arange(self.radius)
    angle_bins = np.arange(-180,180,10)
    dist_index = np.digitize(closest_dist,dis_bin) -1
    theta_index = np.digitize(plane_theta,angle_bins) -1
    rho_index = np.digitize(plane_rho,angle_bins) -1
    return dist_index,theta_index,rho_index

def _state_in_seq(self):
    seq = []
    for plane in self._agent_location:
        m, n = plane
        seq.append(m * self.GRID.shape[1] + n)
    return seq

def render(self):
    rend = self.GRID.copy().astype(dtype = 'U2')
    for plane in range(self.planes):
        rend[self._agent_location[plane][0],self._agent_location[plane][1]] =
→ f'p{plane}'
        rend[self._target_location[plane][0],self.
→_target_location[plane][1]] = f'd{plane}'
    return print(rend)

def reset(self):
    self._agent_location = []
    self._target_location = []
    if self.seed is not None:
        np.random.seed(seed = self.seed)
    for p in range(self.planes):

```

```

        while (pc:= [np.random.randint(self.GRID_DIM[0]),np.random.
↪randint(self.GRID_DIM[1])]) in self._agent_location:
            continue
        self._agent_location.append(pc)
        while (dc:= [np.random.randint(self.GRID_DIM[0]),np.random.
↪randint(self.GRID_DIM[1])]) in self._agent_location + self._target_location:
            continue
        self._target_location.append(dc)
    self._agent_location = np.array(self._agent_location)
    self._target_location = np.array(self._target_location)
    self._count = [0]*self.planes

```

0.4 Random Policy Agent

```

[3]: def random_policy(env,episodes):
    action_space ={
        0: np.array([-1, 0]),
        1: np.array([-1, 1]),
        2: np.array([0, 1]),
        3: np.array([-1 , -1]),
        4: np.array([0 , -1]),
    }
    timestep_reward = []
    cumsum_reward = []
    cumsum_ep= []
    for ep in range(episodes):
        env.reset()
        done = [False]*env.planes
        total_reward = [0]*env.planes
        t=0
        while (not np.array(done).all()) and env._did_collide() and t<1e6:

            action = []
            distance,closest_dist,plane_theta,plane_rh = env._get_info()
            for plane in range(env.planes):
                if closest_dist[plane]< env.radius:
                    action.append(np.random.randint(env.action_space.n))
                else:
                    if np.array_equal(env._target_location[plane],env.
↪_agent_location[plane]):
                        action.append(0)
                    else:
                        x = np.sign(env._target_location[plane] - env.
↪_agent_location[plane])
                        if x[0] in [0,-1]:
                            for act,a in enumerate(action_space.values()):
                                _ = np.array_equal(x,a)

```

```

        if _:
            print(_)
            action.append(act)
            break
        else:
            x = -x
            for act,a in enumerate(action_space.values()):
                _ = np.array_equal(x,a)
                if _:
                    action.append(act)
                    break

            t+=1
            obs,reward,done,info = env.step(action)
            total_reward= np.add(total_reward,reward)
            print(t,ep,end=' ')
            clear_output(wait=True)
            cumsum_reward.append(total_reward/t)
            timestep_reward.append(np.average(total_reward))
            cumsum_ep.append(np.average(total_reward)/(ep+1))
            clear_output()
    return timestep_reward,cumsum_reward,cumsum_ep

```

```

[4]: env = AirTraffic(planes=25,grid_size=[100,100],radius=50,destinations=1)
    episodes = 500
    timestep_reward,cumsum_reward_rand,cumsum_ep_rand = random_policy(env,episodes)

```

```


[75]: # cumsum_plane_avg_rand = np.average(cumsum_reward_rand,axis=1)
    plt.figure(figsize = (10,6))
    plt.title(label=f'Average Score over Time - Random policy ({np.
    ↪average(cumsum_ep_rand[-50:]) :.2f})')
    plt.plot(cumsum_ep_rand)
    plt.xlabel('Epochs')
    plt.ylabel(f'Average Score')

```

```

[75]: Text(0, 0.5, 'Average Score')

```



0.5 Greedy Policy Agent

```
[6]: def greedy_policy(env,episodes):
    action_space ={
        0: np.array([-1, 0]),
        1: np.array([-1, 1]),
        2: np.array([0, 1]),
        3: np.array([-1 , -1]),
        4: np.array([0 , -1]),
    }
    timestep_reward = []
    cumsum_reward = []
    cumsum_ep= []
    max_step=1e6
    for ep in range(episodes):
        env.reset()
        # print(env._get_obs())
        done = [False]*env.planes
        total_reward = [0]*env.planes
        t=0
        while (not np.array(done).all()) and env._did_collide() and t<max_step:

            action = []
            for plane in range(env.planes):
```



```

        x = np.sign(env._target_location[plane] - env.
→_agent_location[plane])
        if x[0] in [0,-1]:
            for act,a in enumerate(action_space.values()):
                _ = np.array_equal(x,a)
                if _:
                    # print(_)
                    action.append(act)
                    break
            else:
                action.append(0)
        else:
            x = -x
            for act,a in enumerate(action_space.values()):
                _ = np.array_equal(x,a)
                if _:
                    action.append(act)
                    break

        t+=1
        obs,reward,done,info = env.step(action)
        total_reward= np.add(total_reward,reward)
        print(t,ep,end=' ')
        clear_output(wait=True)
        cumsum_reward.append(total_reward/t)
        timestep_reward.append(np.average(total_reward))
        cumsum_ep.append(np.average(total_reward)/(ep+1))
        clear_output()
    return timestep_reward,cumsum_reward,cumsum_ep

```

```

[7]: env = AirTraffic(planes=25,grid_size=[100,100],radius=50,destinations=1)
    episodes = 500
    timestep_reward,cumsum_reward_greedy,cumsum_ep_greedy =
→greedy_policy(env,episodes)

```

```


[76]: plt.figure(figsize = (10,6))
    plt.title(label=f'Average Score over Time - Greedy policy ({np.
→average(cumsum_ep_greedy[-50:]) :.2f})')
    plt.plot(cumsum_ep_greedy)
    plt.xlabel('Epochs')
    plt.ylabel(f'Average Score')

```

```

[76]: Text(0, 0.5, 'Average Score')

```



0.6 Multi-Agent SARSA

```
[9]: def ep_greedy(env,Q,count_state,epsilon=0.1):
    dist_index,theta_index,rho_index = env._state_to_index()
    action = []
    for plane in range(env.planes):
        if np.random.random()>epsilon:
            #
            ↪x=(count_state[dist_index[plane],theta_index[plane],rho_index[plane],:]!=0).
            ↪all()
            if True :
                action.append(np.
            ↪argmax(Q[dist_index[plane],theta_index[plane],rho_index[plane],:]))
            else:
                axn = np.
            ↪where(Q[dist_index[plane],theta_index[plane],rho_index[plane],:]==0)[0]
                action.append(axn[0])
            else:
                action.append(np.random.randint(env.action_space.n))
    return action

def Sarsa(env,alpha, gamma, epsilon, episodes,temperature=None):
    cumsum_reward = []
```

```

count_state = np.zeros_like(Q)
Q_temp = np.zeros_like(Q)
for ep in range(episodes):
    env.reset()
    done = [False]*env.planes
    total_reward = [0]*env.planes
    curr_dist_index,curr_theta_index,curr_rho_index = env._state_to_index()
    curr_a = ep_greedy(env,Q,epsilon=epsilon,count_state=count_state)
    t = 0
    while (not np.array(done).all()) and env._did_collide() and t<1.2e4:
        for plane in range(env.planes):
            ↪count_state[curr_dist_index[plane],curr_theta_index[plane],curr_rho_index[plane],curr_a[plane]]
            obs, reward, done, info = env.step(curr_a)
            nxt_dist_id,nxt_theta_id,nxt_rho_id = env._state_to_index()
            total_reward= np.add(total_reward,reward)
            next_act = ep_greedy(env,Q,epsilon=epsilon,count_state=count_state)
            t+=1

            for plane in range(env.planes):
                ↪
                ↪Q[curr_dist_index[plane],curr_theta_index[plane],curr_rho_index[plane],curr_a[plane]]
                ↪+= alpha * ( reward[plane] + (gamma *
                ↪Q[nxt_dist_id[plane],nxt_theta_id[plane],nxt_rho_id[plane], next_act[plane]]
                ↪)
                ↪
                ↪Q[curr_dist_index[plane],curr_theta_index[plane],curr_rho_index[plane],curr_a[plane]]
                ↪)

                curr_dist_index,curr_theta_index,curr_rho_index =
                ↪nxt_dist_id,nxt_theta_id,nxt_rho_id
                curr_a = next_act
                print(t,ep,end=' ')
                clear_output(wait=True)
                cumsum_reward.append(total_reward/t)
                timestep_reward.append(np.average(total_reward))
                cumsum_ep.append(np.average(total_reward)/(ep+1))
                if temperature is not None and ep%temperature==0:
                    epsilon = 0.9*epsilon
                    print(epsilon)

clear_output()
return timestep_reward,cumsum_reward, cumsum_ep


```

0.6.1 Constant Temperature SARSA

```
[10]: env = AirTraffic(planes=25,grid_size=[100,100],radius=50,destinations=1)
n_radius,n_theta, n_rho,n_actions = env.radius, 36,36, env.action_space.n
Q = np.zeros((n_radius,n_theta, n_rho,n_actions))
alpha = 0.5
gamma = 0.9
epsilon = 0.1
episodes = 500
cumsum_ep= []
timestep_reward = []
timestep_reward,cumsum_reward,cumsum_ep_sarsa = Sarsa(env,alpha, gamma,
↪epsilon, episodes)
```

```
[77]: # cumsum_plane_avg_rand = np.average(cumsum_reward_rand,axis=1)
plt.figure(figsize = (10,6))
plt.title(label=f'Average Score over Time - SARSA policy ({np.
↪average(cumsum_ep_sarsa[-50:]) :.2f})')
plt.plot(cumsum_ep_sarsa,label = r'($\epsilon$ = ' + f'{epsilon}','+r'$\alpha$ = '
↪' + f'{alpha}', ' + r'$\gamma$ = ' + f'{gamma})')
plt.xlabel('Epochs')
plt.ylabel(f'Average Score')
plt.legend()
```

[77]: <matplotlib.legend.Legend at 0x1c7da9c42e0>




0.6.2 Temperature Controlled Exploration

```
[12]: env = AirTraffic(planes=25,grid_size=[100,100],radius=50,destinations=1)
n_radius,n_theta, n_rho,n_actions = env.radius, 36,36, env.action_space.n
Q = np.zeros((n_radius,n_theta, n_rho,n_actions))
alpha = 0.5
gamma = 0.9
epsilon = 0.5
episodes = 500
cumsum_ep= []
timestep_reward = []
temperature = 10
timestep_reward,cumsum_reward,cumsum_sarsa_temp = Sarsa(env,alpha, gamma,
↪epsilon, episodes,temperature)

[78]: plt.figure(figsize = (10,6))
plt.title(label=f'Average Score over Time - Multi Agent SARSA ({np.
↪average(cumsum_sarsa_temp[-50:]) :.2f})')
plt.plot(cumsum_sarsa_temp,label = r'$\alpha$ = ' +f'{alpha}, '+ r'$\gamma$ = '
↪+f'{gamma}, '+r'$T$ = '+f'{1-1/temperature}')
plt.xlabel('Epochs')
plt.ylabel(f'Average Score')
plt.legend()
```

[78]: <matplotlib.legend.Legend at 0x1c7dab0e310>

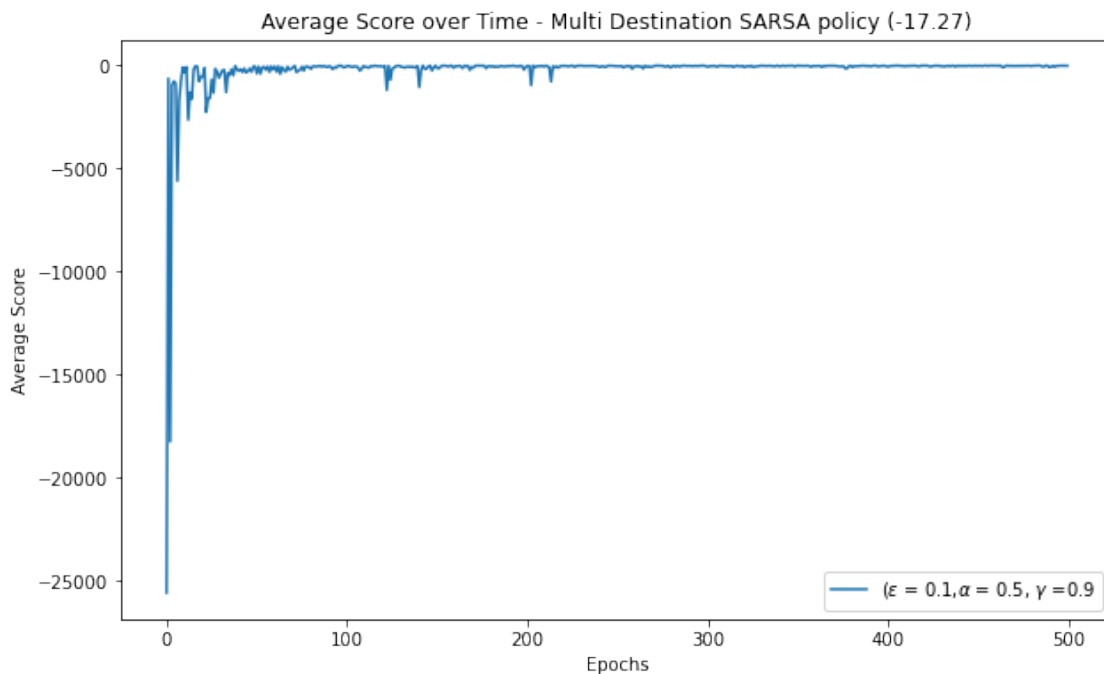


0.7 Multi Agent - Multi Destination SARSA

```
[14]: env = AirTraffic(planes=25,grid_size=[100,100],radius=50)
n_radius,n_theta, n_rho,n_actions = env.radius, 36,36, env.action_space.n
Q = np.zeros((n_radius,n_theta, n_rho,n_actions))
alpha = 0.5
gamma = 0.9
epsilon = 0.1
episodes = 500
cumsum_ep= []
timestep_reward = []
timestep_reward,cumsum_reward_multi,cumsum_sarsa_multi = Sarsa(env,alpha,
↪gamma, epsilon, episodes)
```

```
[79]: # cumsum_plane_avg_rand = np.average(cumsum_reward_rand,axis=1)
plt.figure(figsize = (10,6))
plt.title(label=f'Average Score over Time - Multi Destination SARSA policy ({np.
↪average(cumsum_sarsa_multi[-50:]) :.2f})')
plt.plot(cumsum_sarsa_multi,label = r'($\epsilon$ = ' + f'{epsilon}','+r'$\alpha$ = '
↪= ' + f'{alpha}', '+ r'$\gamma$ = ' + f'{gamma}')
plt.xlabel('Epochs')
plt.ylabel(f'Average Score')
plt.legend()
```

[79]: <matplotlib.legend.Legend at 0x1c7dbca3df0>



0.8 Greedy + SARSA

```
[16]: def Sarsa_policy(env,Q,epsilon,count_state=None):
    action_space = {
        0: np.array([-1, 0]),
        1: np.array([-1, 1]),
        2: np.array([0, 1]),
        3: np.array([-1, -1]),
        4: np.array([0, -1]),
    }

    action = []
    Plane = []
    distance,closest_dist,plane_theta,plane_rh = env._get_info()
    dist_index,theta_index,rho_index = env._state_to_index()
    # print(closest_dist)
    for plane in range(env.planes):
        if closest_dist[plane]< env.radius:
            if np.random.random()>epsilon:
                action.append(np.
→argmax(Q[dist_index[plane],theta_index[plane],rho_index[plane],:]))
            else:
                action.append(np.random.randint(env.action_space.n))
                Plane.append(plane)
            else:
                if np.array_equal(env._target_location[plane],env.
→_agent_location[plane]):
                    action.append(0)
                else:
                    x = np.sign(env._target_location[plane] - env.
→_agent_location[plane])
                    if x[0] in [0,-1]:
                        for act,a in enumerate(action_space.values()):
                            _ = np.array_equal(x,a)
                            if _:
                                # print(_)
                                action.append(act)
                                break
                    else:
                        x = -x
                        for act,a in enumerate(action_space.values()):
                            _ = np.array_equal(x,a)
                            if _:
                                action.append(act)
                                break
    # print(Plane)
    return action, Plane
```

```

def greedy_Sarsa(env,alpha, gamma, epsilon, episodes,temperature=None):

    cumsum_reward = []
    count_state = np.zeros_like(Q)
    Q_temp = np.zeros_like(Q)

    for ep in range(episodes):
        env.reset()
        done = [False]*env.planes
        total_reward = [0]*env.planes
        curr_dist_index,curr_theta_index,curr_rho_index = env._state_to_index()
        curr_a,curr_Planes = _
        Sarsa_policy(env,Q,epsilon=epsilon,count_state=count_state)
        t = 0
        while (not np.array(done).all()) and env._did_collide() and t<1.2e4:
            obs, reward, done, info = env.step(curr_a)
            nxt_dist_id,nxt_theta_id,nxt_rho_id = env._state_to_index()
            total_reward= np.add(total_reward,reward)
            next_act,nxt_Planes = Sarsa_policy(env,Q,epsilon,count_state=None)
            t+=1
            if len(curr_Planes) != 0:
                for plane in curr_Planes:
                    _
                    Q[curr_dist_index[plane],curr_theta_index[plane],curr_rho_index[plane],curr_a[plane]]_
                    += alpha * ( reward[plane] + (gamma *_
                    Q[nxt_dist_id[plane],nxt_theta_id[plane],nxt_rho_id[plane], next_act[plane]])
                    _
                    _
                    Q[curr_dist_index[plane],curr_theta_index[plane],curr_rho_index[plane],curr_a[plane]])
                    curr_dist_index,curr_theta_index,curr_rho_index =_
                    nxt_dist_id,nxt_theta_id,nxt_rho_id
                    curr_a = next_act
                    curr_Planes = nxt_Planes
                    print(t,ep,end=' ')
                    clear_output(wait=True)
#         env.render()
        cumsum_reward.append(total_reward/t)
        timestep_reward.append(np.average(total_reward))
        cumsum_ep.append(np.average(total_reward)/(ep+1))
        if ep%25 ==0:
            Q_temp = Q.copy()
            if temperature is not None and ep%temperature==0:
                epsilon = 0.9*epsilon
                print(epsilon)

    clear_output()

```

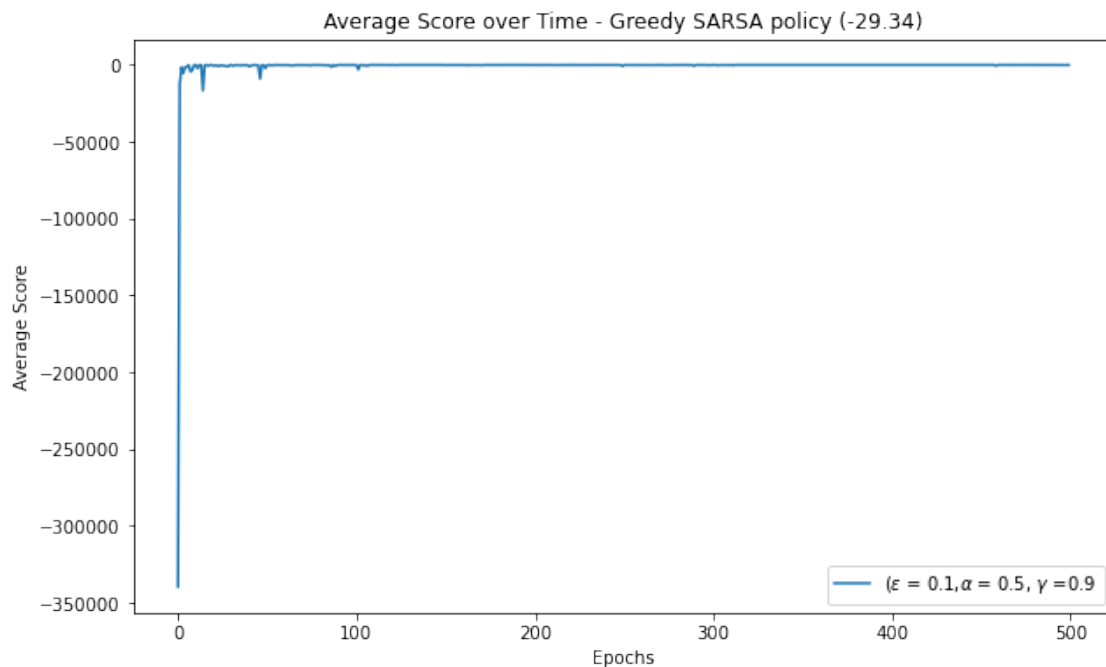


```
return timestep_reward,cumsum_reward,cumsum_ep
```

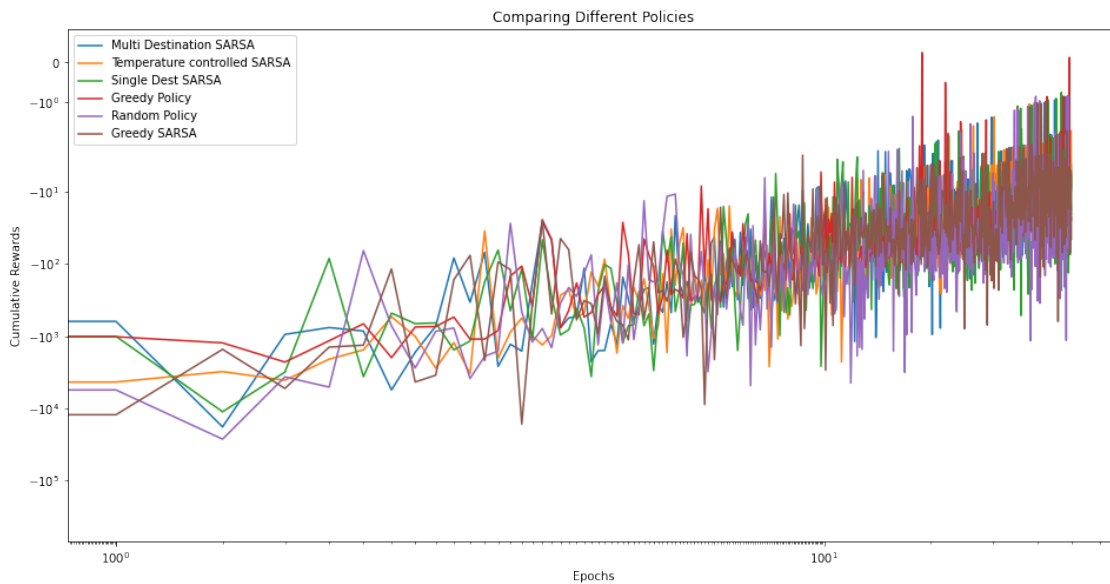
```
[17]: env = AirTraffic(planes=25,grid_size=[100,100],radius=50,destinations=1)
n_radius,n_theta, n_rho,n_actions = env.radius, 36,36, env.action_space.n
Q = np.zeros((n_radius,n_theta, n_rho,n_actions))
alpha = 0.5
gamma = 0.9
epsilon = 0.1
episodes = 500
cumsum_ep= []
timestep_reward = []
timestep_reward,cumsum_reward_greedy_sarsa,cumsum_greedy_sarsa =
↳greedy_Sarsa(env,alpha, gamma, epsilon, episodes)
```

```
[80]: plt.figure(figsize = (10,6))
plt.title(label=f'Average Score over Time - Greedy SARSA policy ({np.
↳average(cumsum_greedy_sarsa[-50:]) :.2f}))')
plt.plot(cumsum_greedy_sarsa,label = r'($\epsilon$ = '+
↳f'{epsilon}','+r'$\alpha$ = ' +f'{alpha}', '+ r'$\gamma$ = ' +f'{gamma}'))
plt.xlabel('Epochs')
plt.ylabel(f'Average Score')
plt.legend()
```

[80]: <matplotlib.legend.Legend at 0x1c7dbd7b910>



```
[71]: plt.figure(figsize = (16,8))
plt.plot(cumsum_sarsa_multi,label = 'Multi Destination SARSA')
plt.plot(cumsum_sarsa_temp,label = 'Temperature controlled SARSA')
plt.plot(cumsum_ep_sarsa, label = 'Single Dest SARSA')
plt.plot(cumsum_ep_greedy,label = 'Greedy Policy')
plt.plot(cumsum_ep_rand, label = 'Random Policy')
plt.plot(cumsum_greedy_sarsa,label = 'Greedy SARSA')
plt.legend()
plt.xlabel('Epochs')
plt.ylabel('Cumulative Rewards')
plt.title('Comparing Different Policies')
plt.yscale('symlog')
plt.semilogx(base = 100)
plt.show()
```



```
[ ]:
```